

Hi Ivan,

Glad we are aligned. Formally closing the first stage and pivoting our focus entirely to the core quality of the model (text and phonetics) is absolutely the right technical decision given the hardware constraints we identified.

1. Transferring Funds from Stage 1 to Stage 2 The cleanest way to handle the budget transfer depends on how the current contract is structured:

If on a platform like Freelancer: We can simply edit the active Milestone 1, rename it to "Milestone 2: Multilingual Phonetic Fine-Tuning," and adjust the deliverables/amount to reflect our new Phase 2 scope.

If working directly: We can cancel the remaining deliverables of Invoice/Stage 1, apply the remaining balance as a direct credit to Stage 2, and I will issue a revised Statement of Work (SoW) for the new milestone. Let me know which method you prefer.

2. Advisory: Text/Phonetics vs. Personality/Manner of Speech I strongly advise that we separate base language/phonetics training from "personality and manner of speech." Here is why:

Step 1: Text & Phonetics (Do this NOW): We train a highly efficient, low-rank LoRA adapter ($r=16$) on Moshi's text projection layers and audio codebooks. This teaches the model how to speak Russian/Uzbek with correct pronunciation and grammar. Because the adapter is lightweight, it adds effectively zero latency to the inference pipeline, preserving the speed you need.

Step 2: Personality & Manner of Speech (Do this LATER / Dynamically): Instead of permanently baking a single personality into the model weights (which risks overfitting and latency), we control personality dynamically. I have already built an isolated microservice ("PersonaPlex") that injects dynamic text prompts directly into the Moshi session in real-time. We can define the character, emotional state, and mannerisms via text context.

This approach gives you maximum flexibility: a phonetically perfect baseline model that can instantly switch "personalities" on the fly without requiring expensive retraining.

3. Proposed Milestone 2 Agreement (For Your Review)

Objective: Train and deploy a lightweight QLoRA adapter for Moshi targeting bilingual (Russian/Uzbek) text and audio codebooks to achieve native phonetic pronunciation with minimal latency overhead.

Deliverable 1: Training Configuration & Setup. Audit the Moshi tokenizer for Cyrillic/Latin compatibility and configure the training scripts (Rank, Alpha, Dropout) for latency-aware response shaping.

Deliverable 2: Phonetic Adapter Training. Execute the fine-tuning on the provided bilingual datasets.

Deliverable 3: Evaluation & Optimization. Deploy the trained adapter onto the H100 server and evaluate it for native fluency, grammar accuracy, and stable real-time streaming performance (ensuring no degradation in latency).

Deliverable 4: PersonaPlex Integration. Hook the fine-tuned model into the new PersonaPlex microservice to prove that we can dynamically alter the model's emotional state and personality using real-time text prompting over the trained phonetics.

Take your time preparing the datasets. Once you are ready, we will discuss the exact parameter settings, finalize the milestone transfer, and begin the training run.

Best,

Novatech